

# Query processing

a
2
3

"select a from X natural join Y where c = 3;"

X

a	b
1	aaa
2	bbb
3	ccc


Y

b	c
aaa	1
bbb	2
bbb	3
ccc	3
ddd	4


# Compilation

SQL { "select a from X natural join Y where c = 3;"

Relational  
algebra {



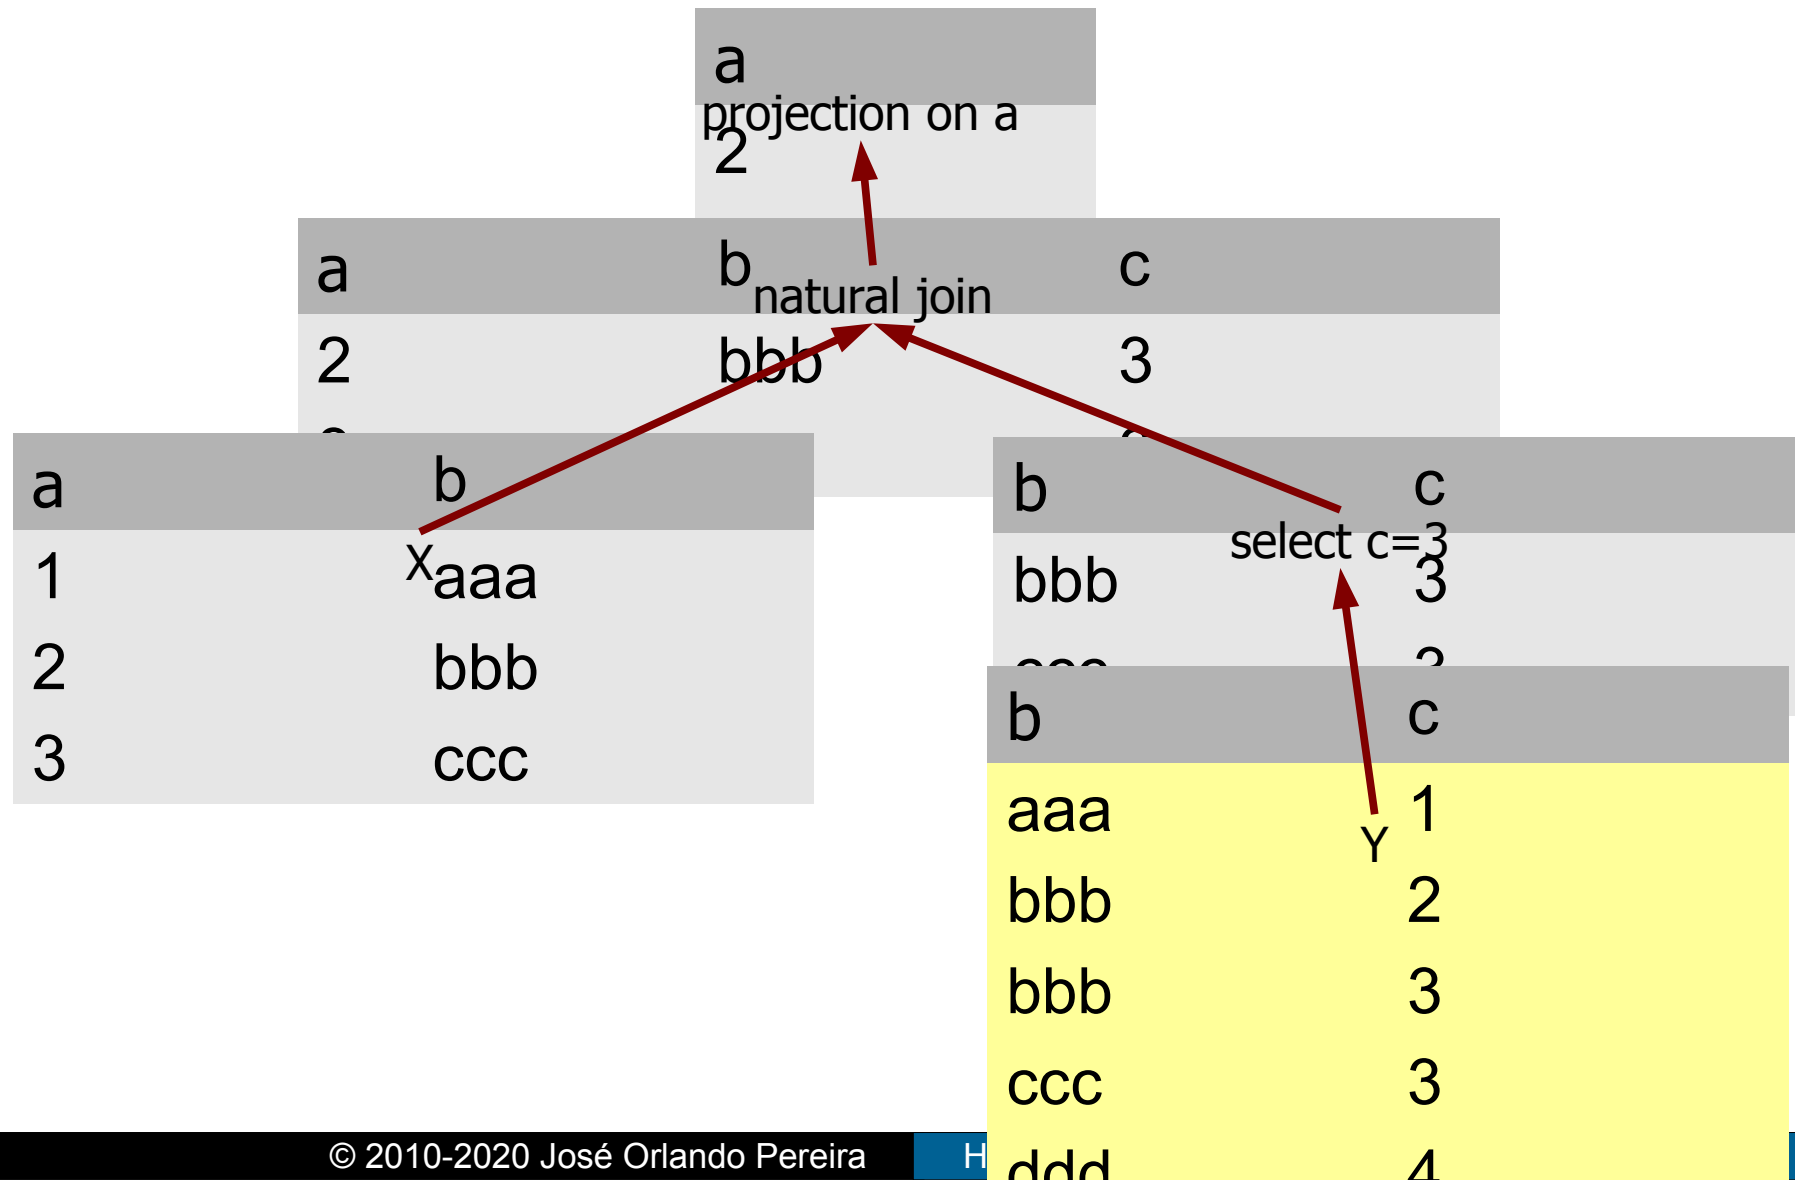
# Logical execution



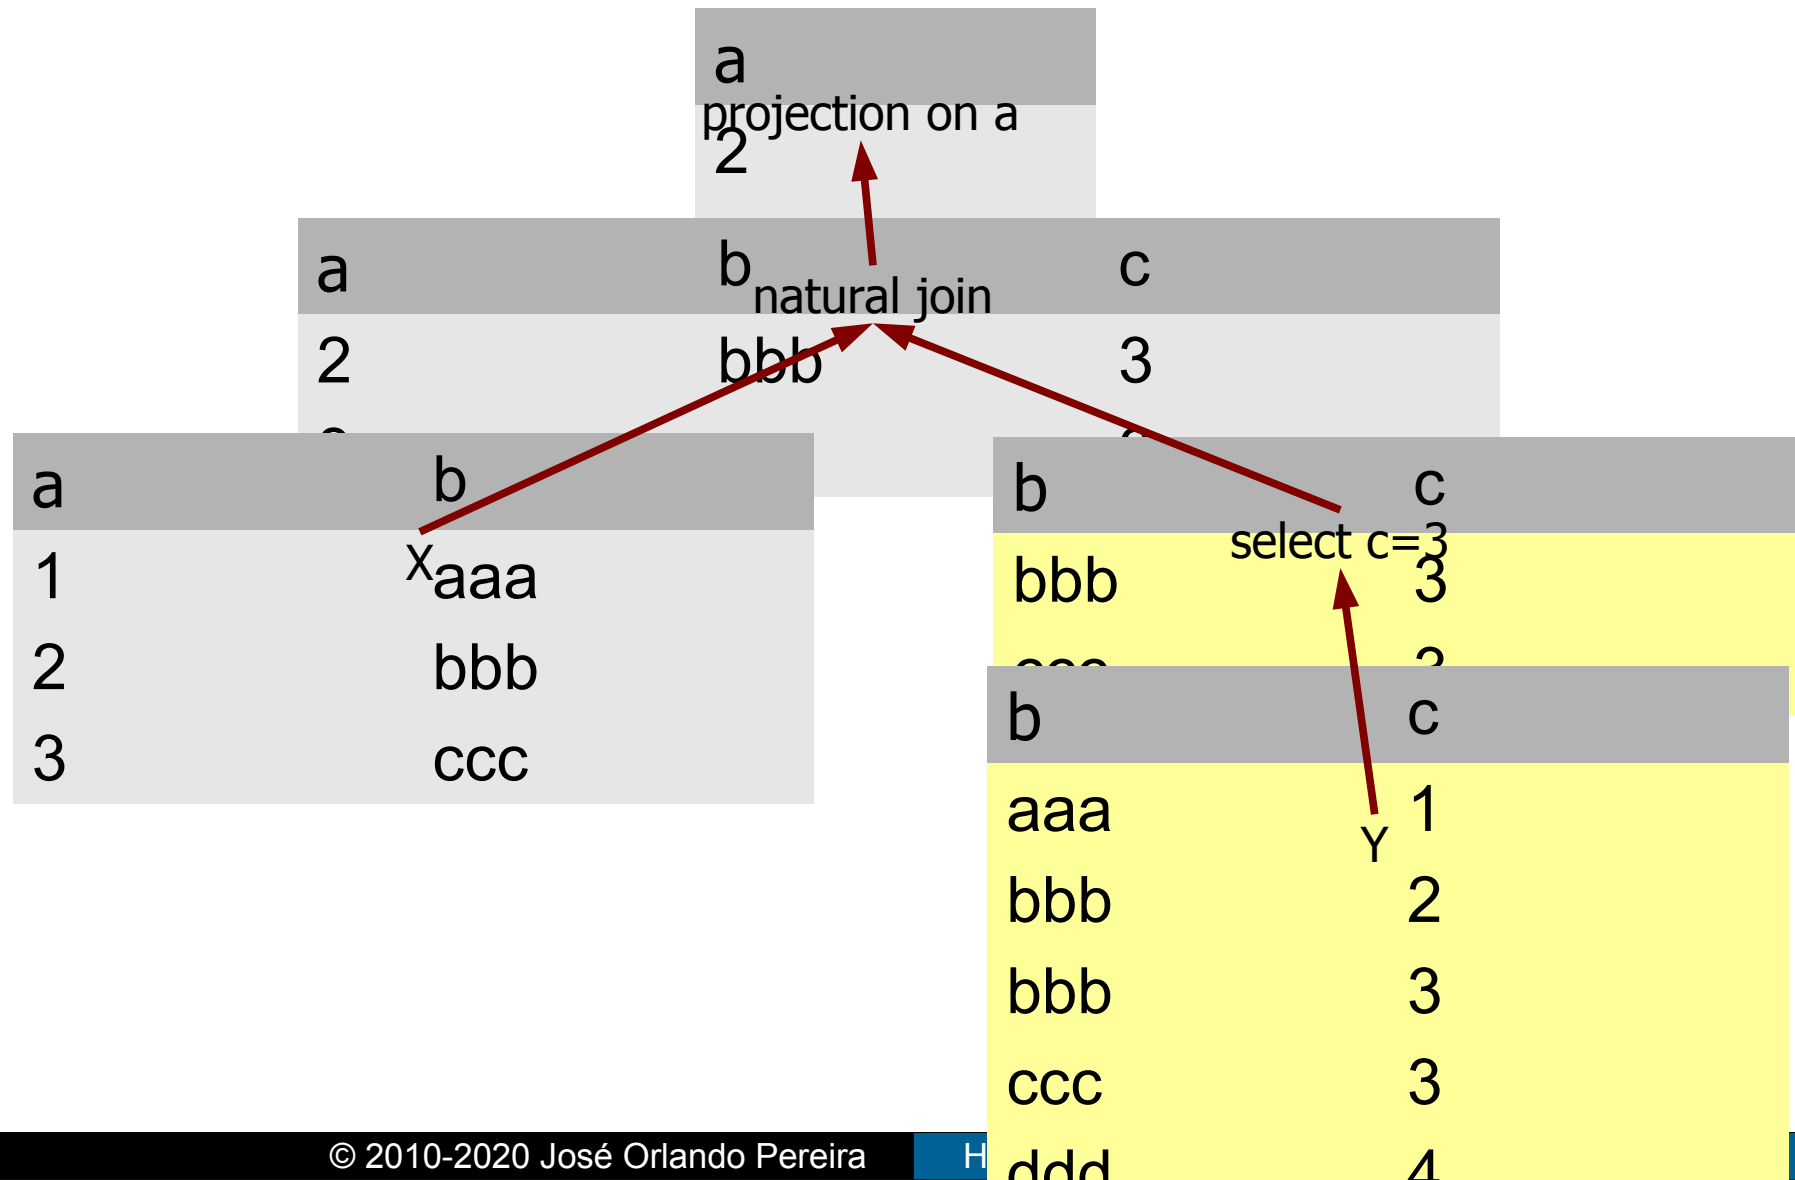
# Materialization

- Each operator is a function:
  - Returns a relation
  - Parameters are other relations (possibly, returned from operators)
- Computation order:
  - From leaves to root


# Execution with materialization



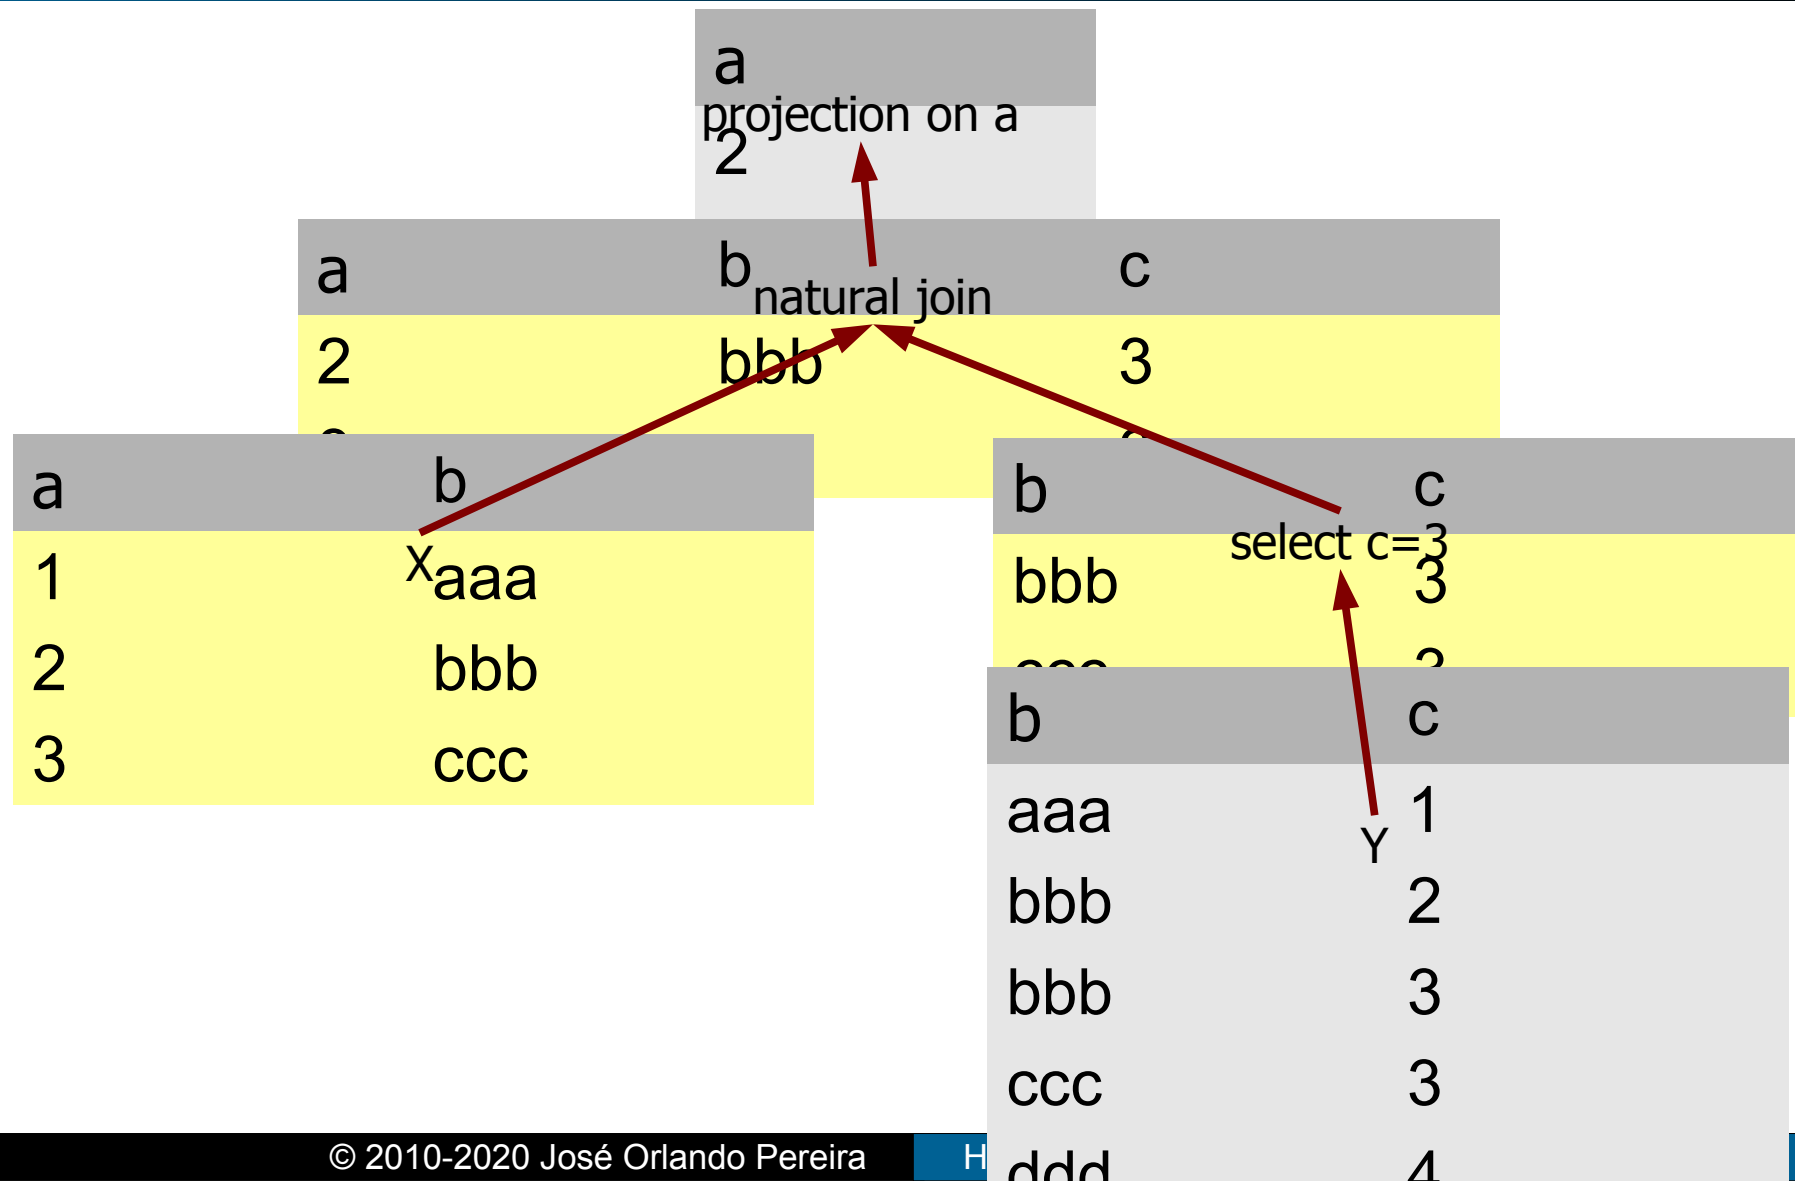
# Execution with materialization



# Execution with materialization

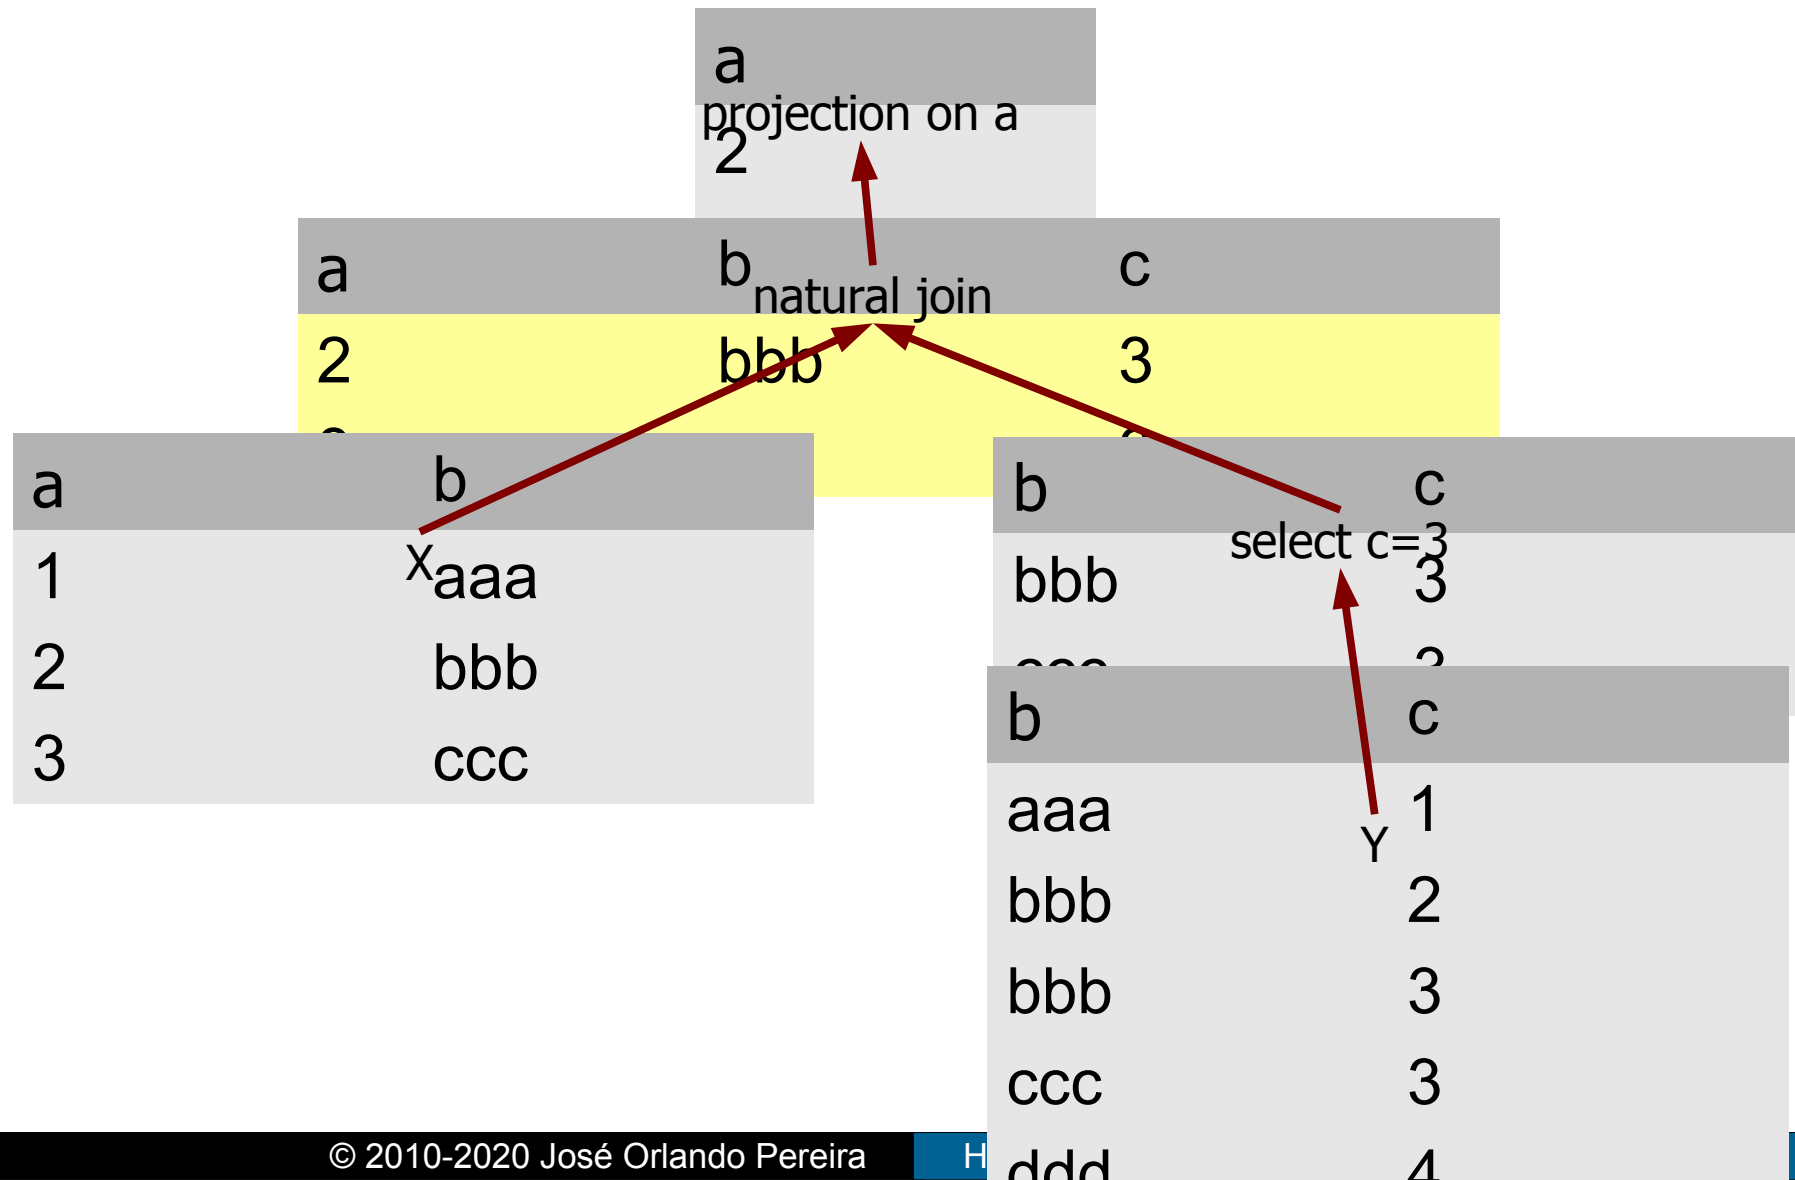


# Execution with materialization

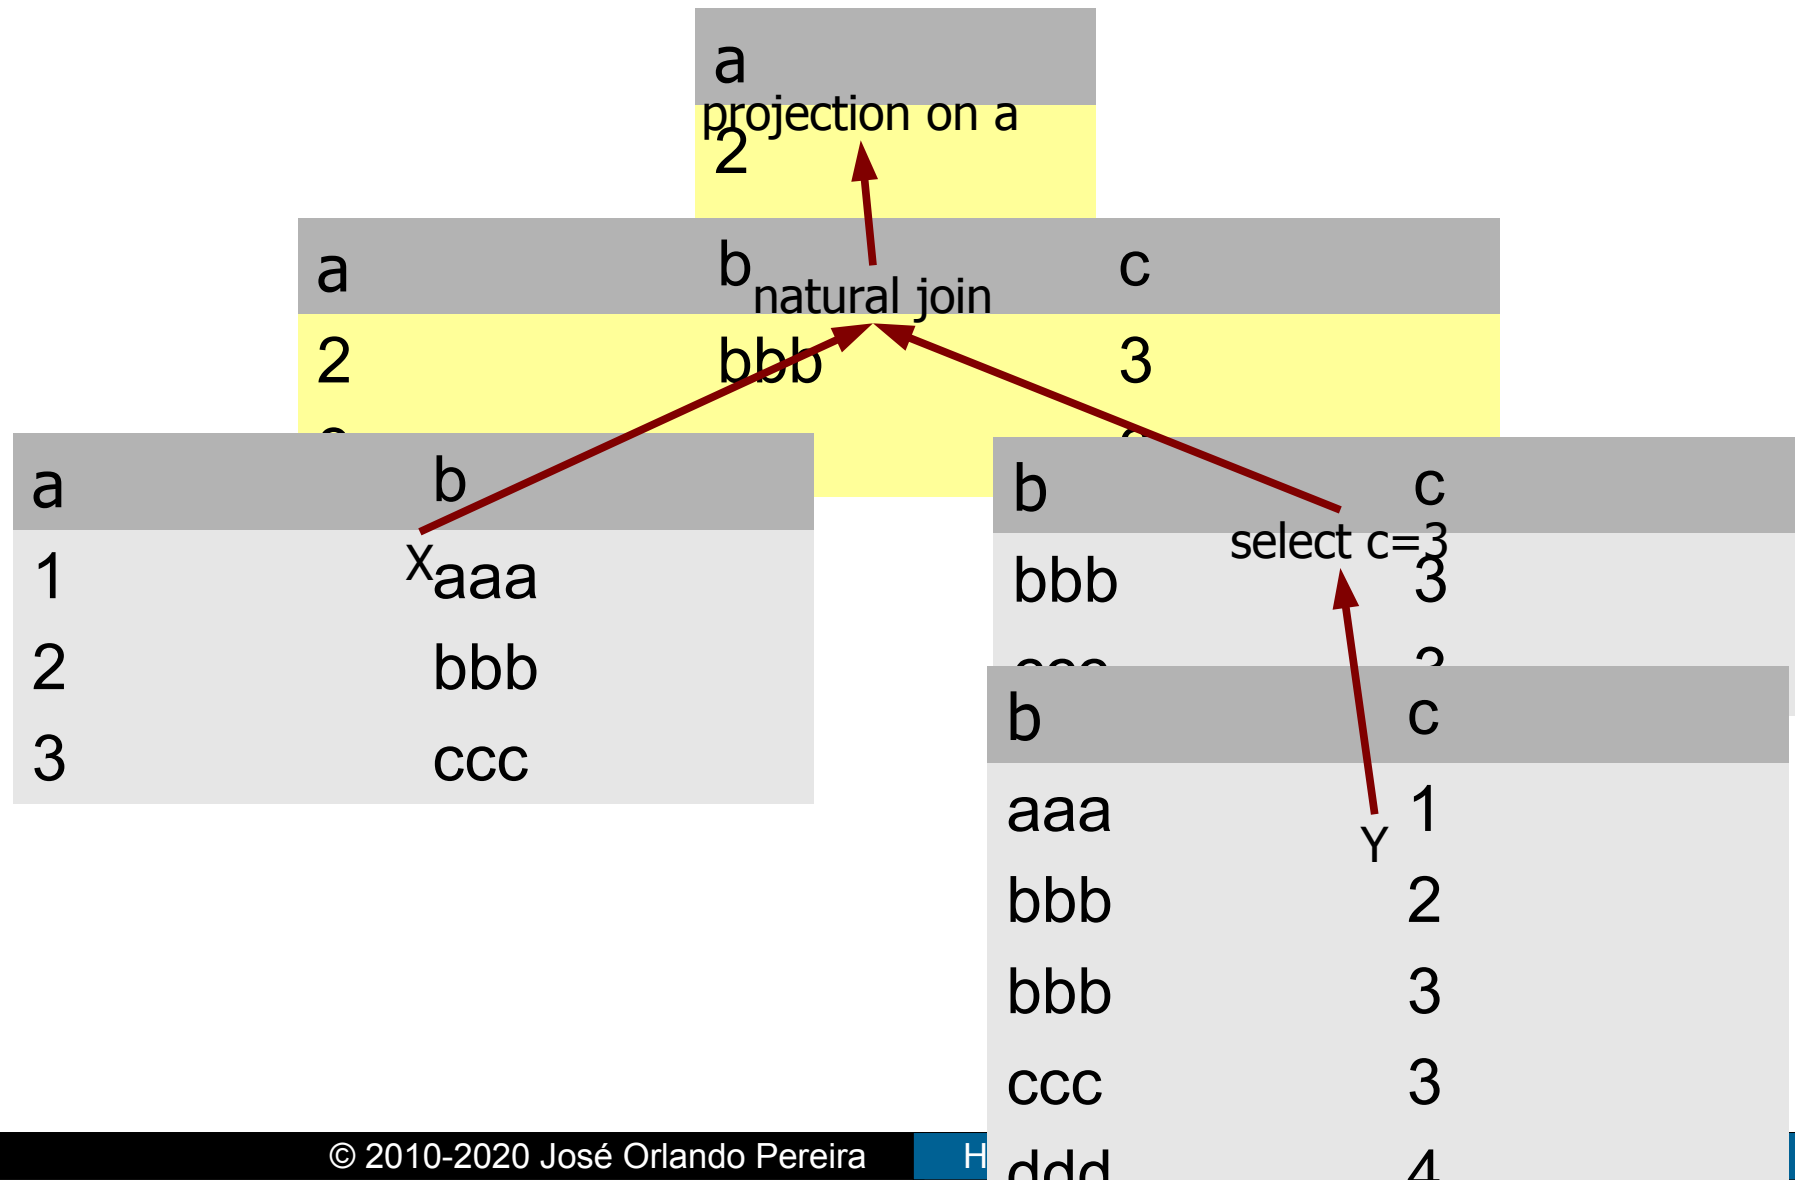





# Execution with materialization



# Execution with materialization




# Execution with materialization



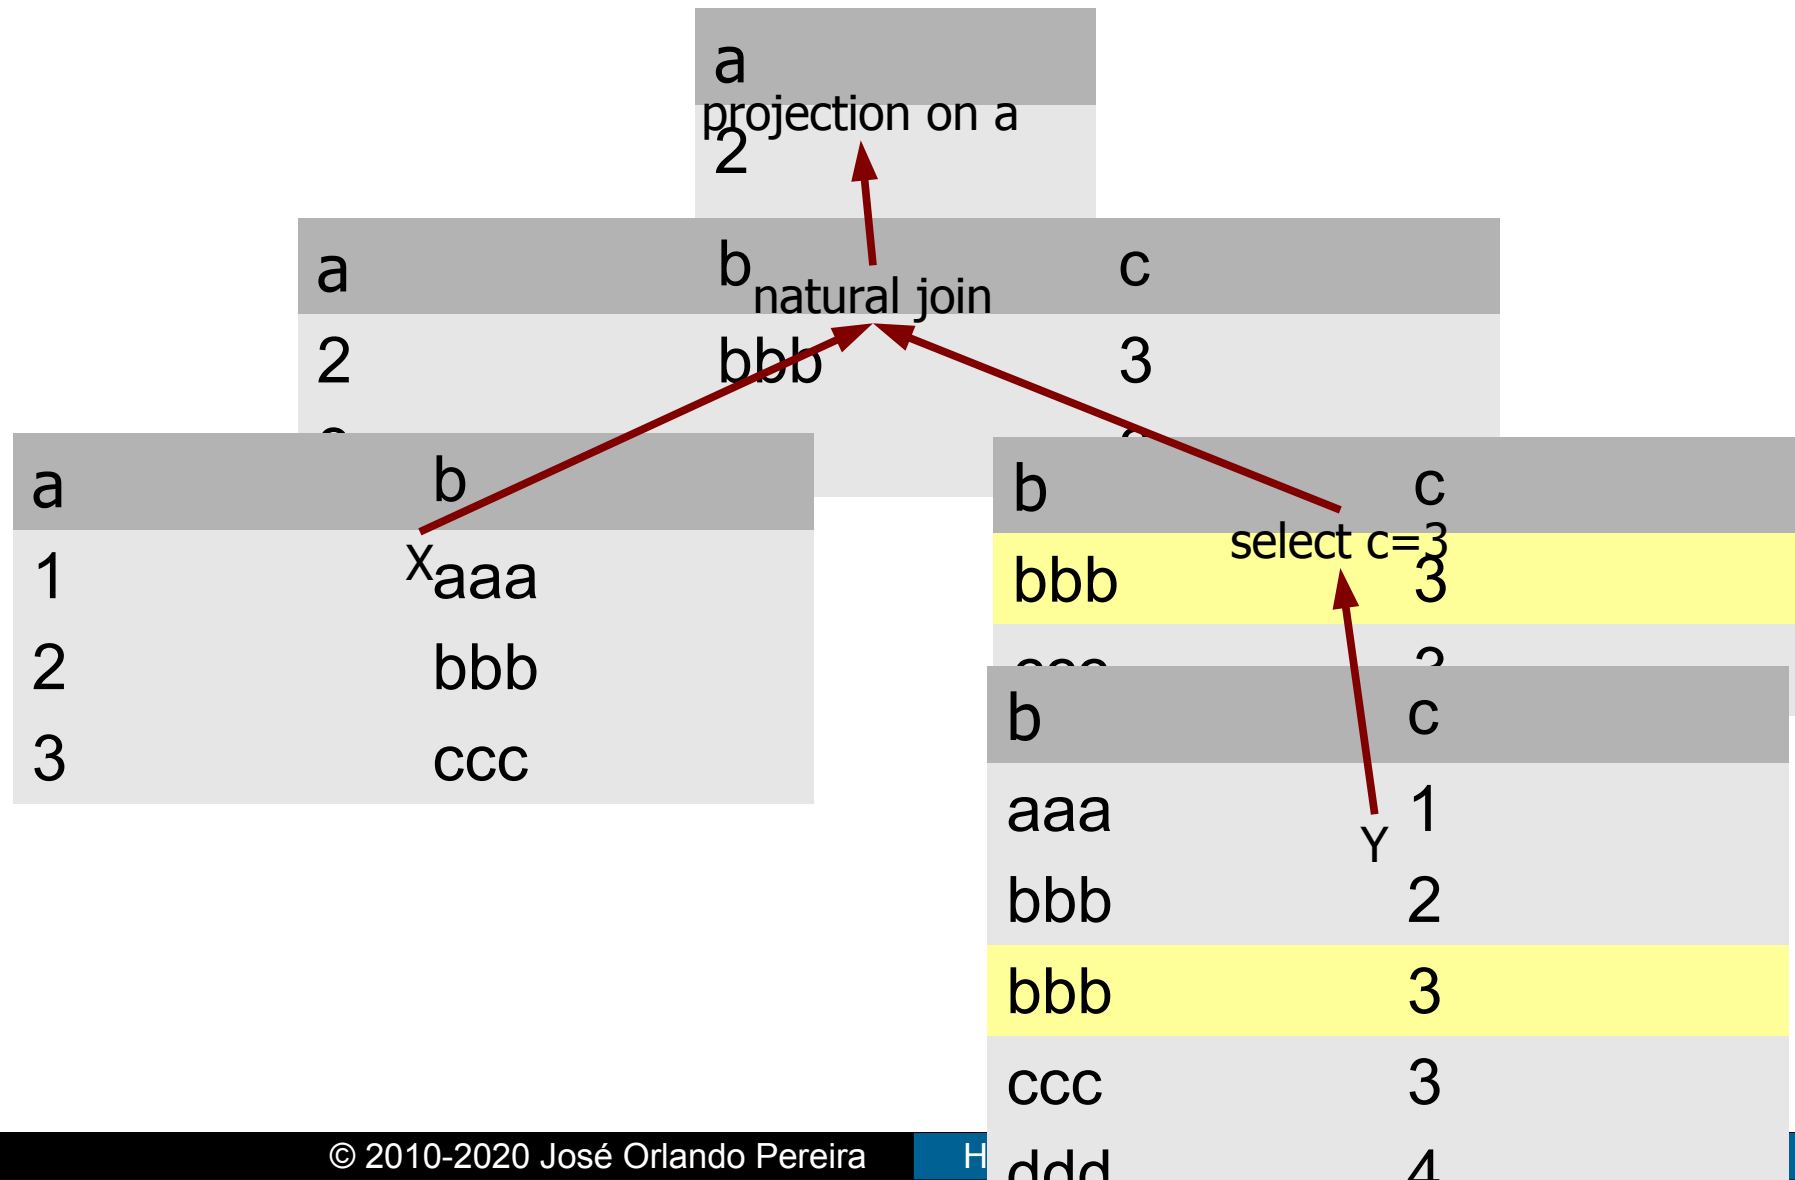
# Iteration

- Each operator is an object:
  - Interface similar to `java.util.Iterator`:
    - `open()` - get ready to return first record
    - `next()` - return next record
    - `close()` - no more records required
  - Constructor parameters:
    - Other operator objects
- Computation order:
  - From leaves to root, for each record

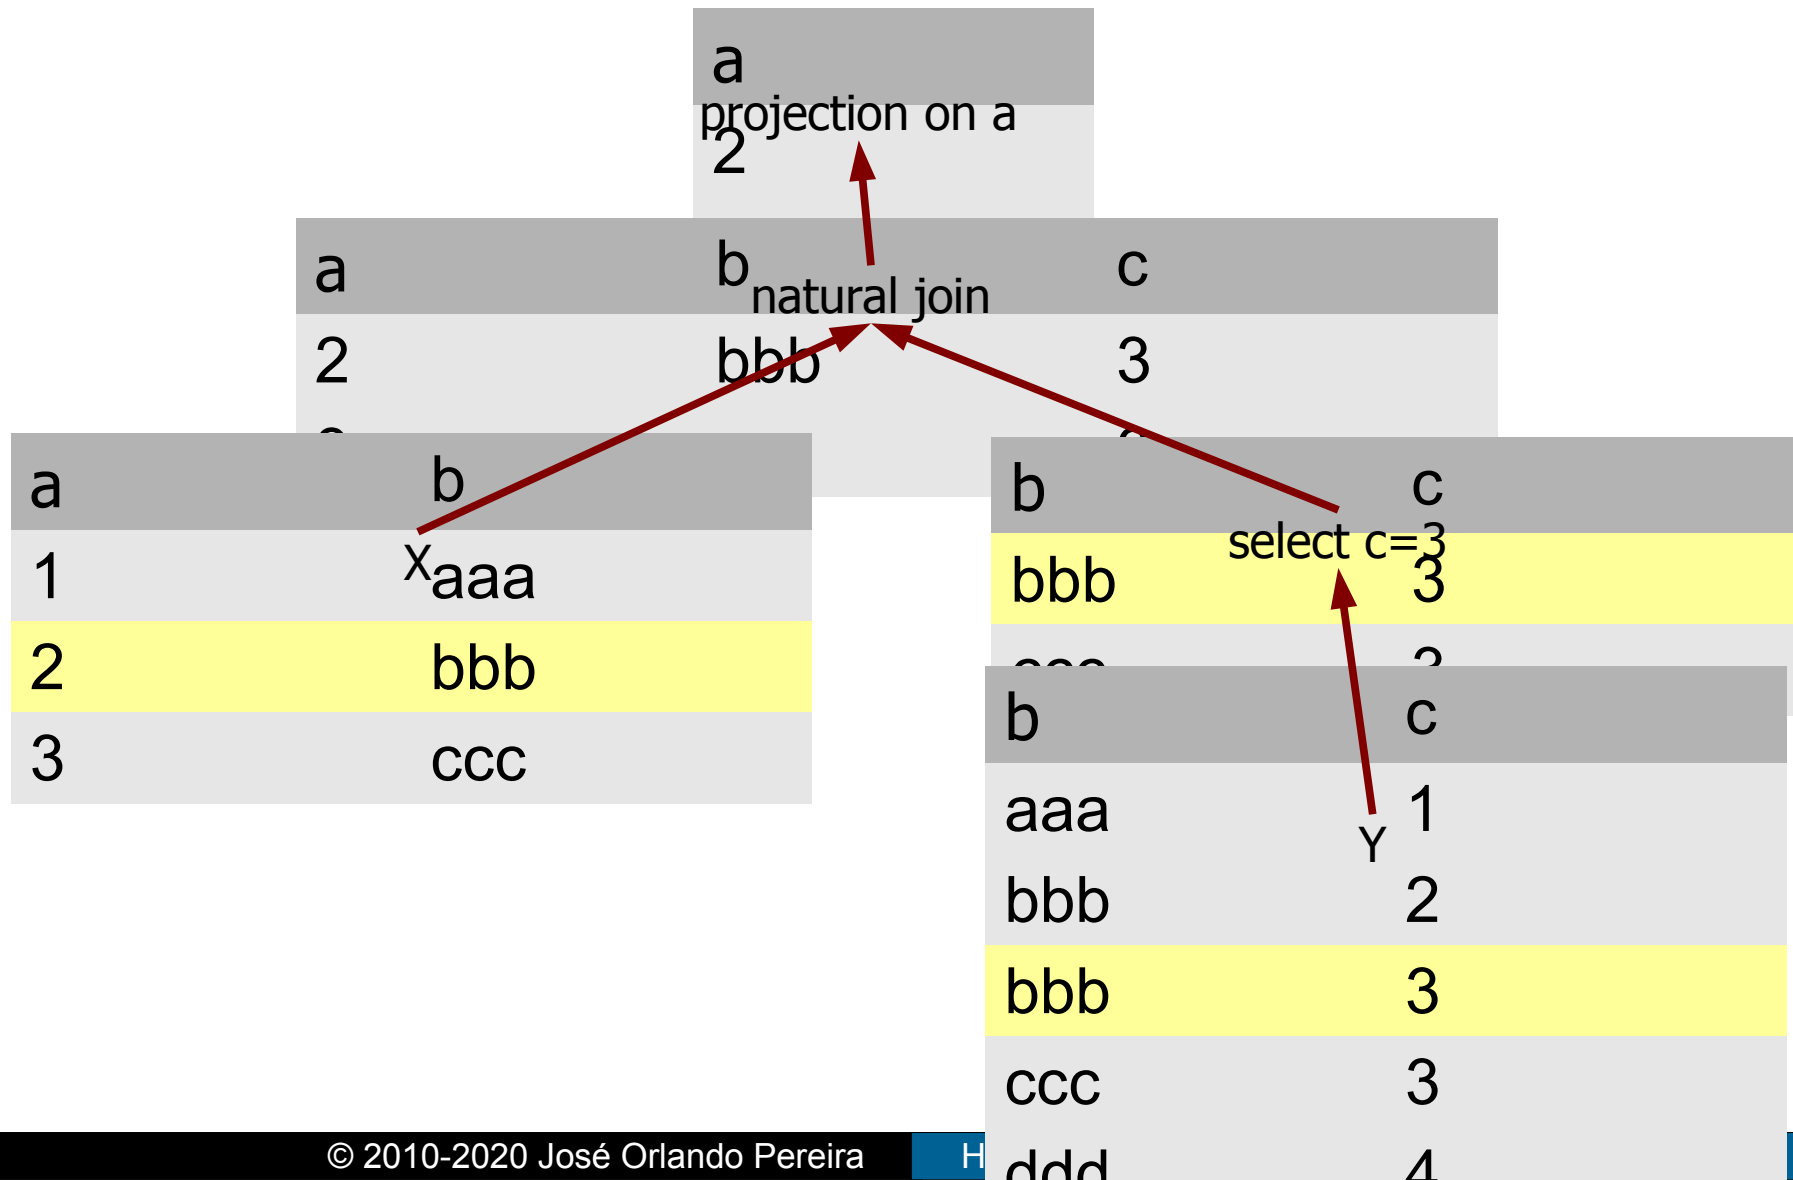
# Execution with iteration



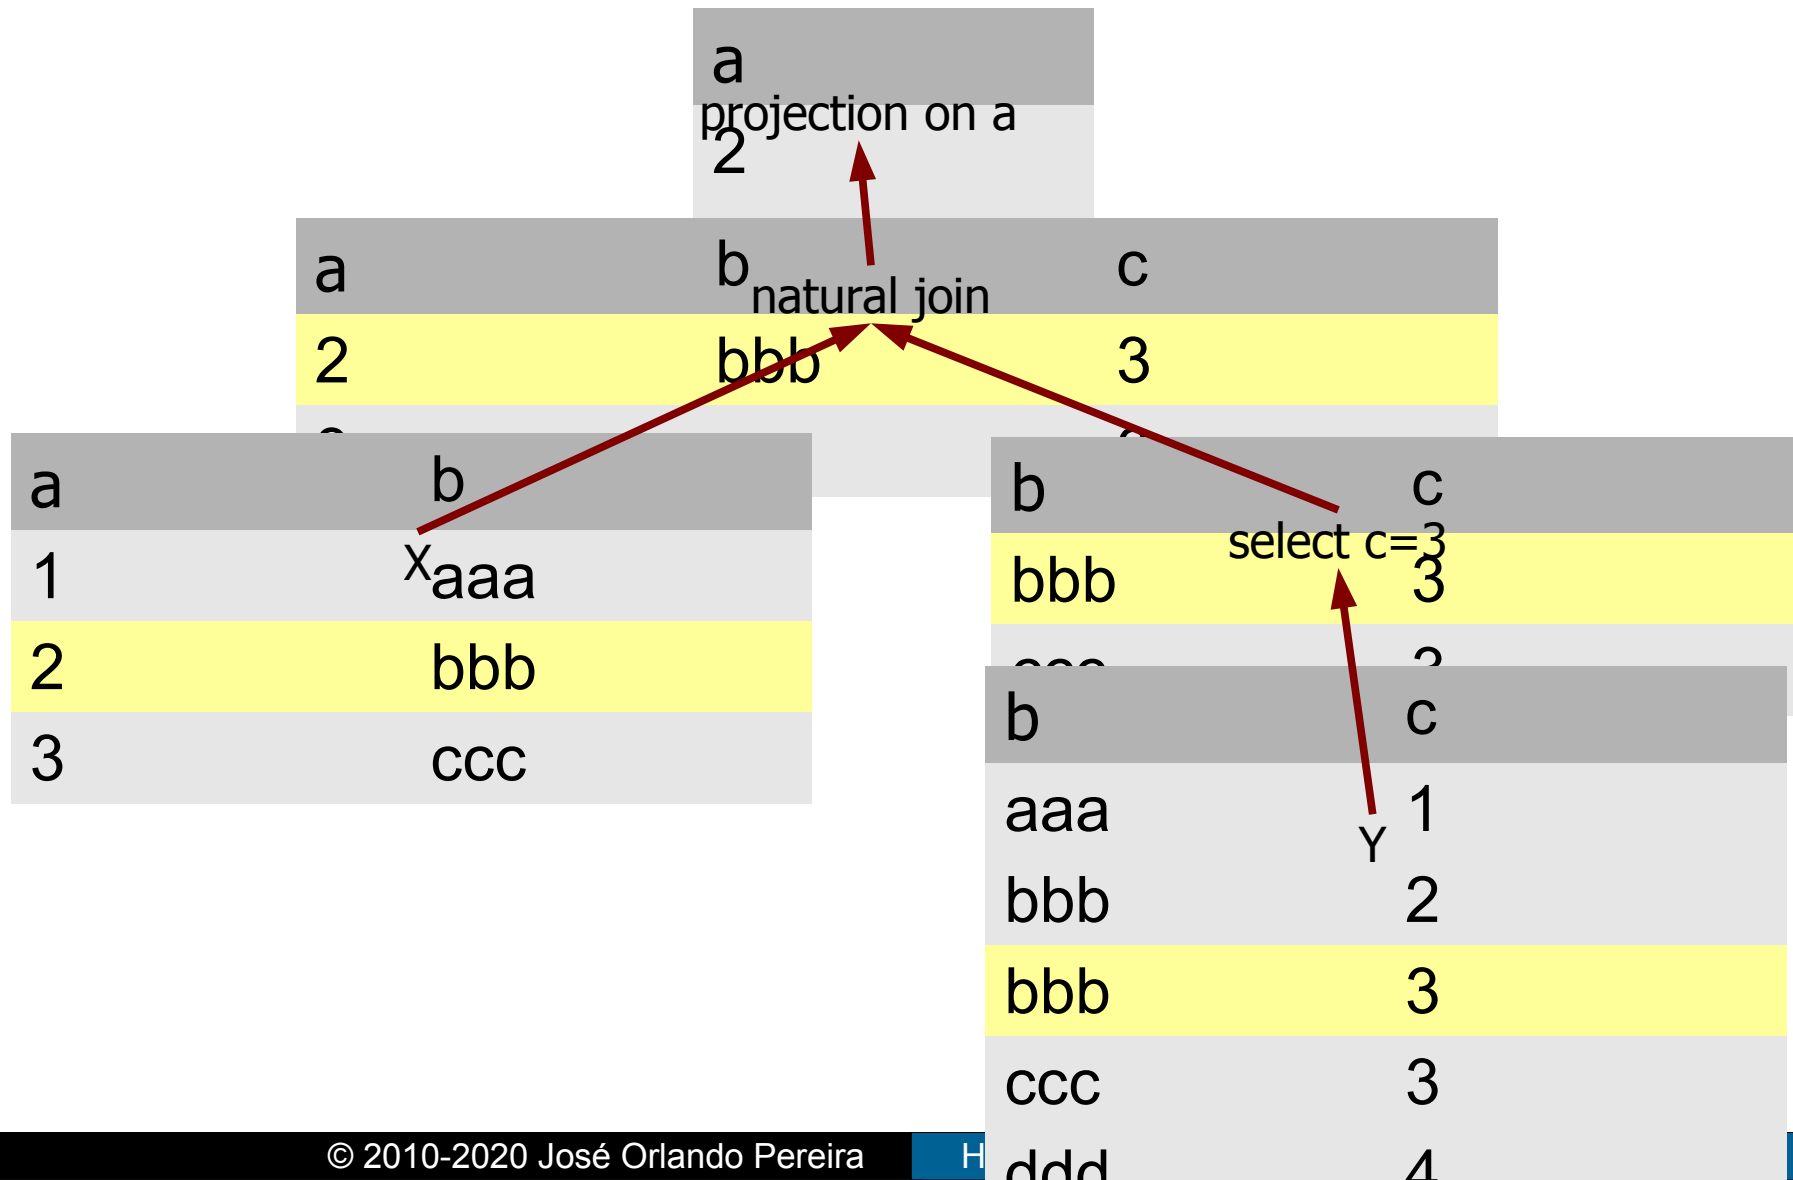
# Execution with iteration



# Execution with iteration

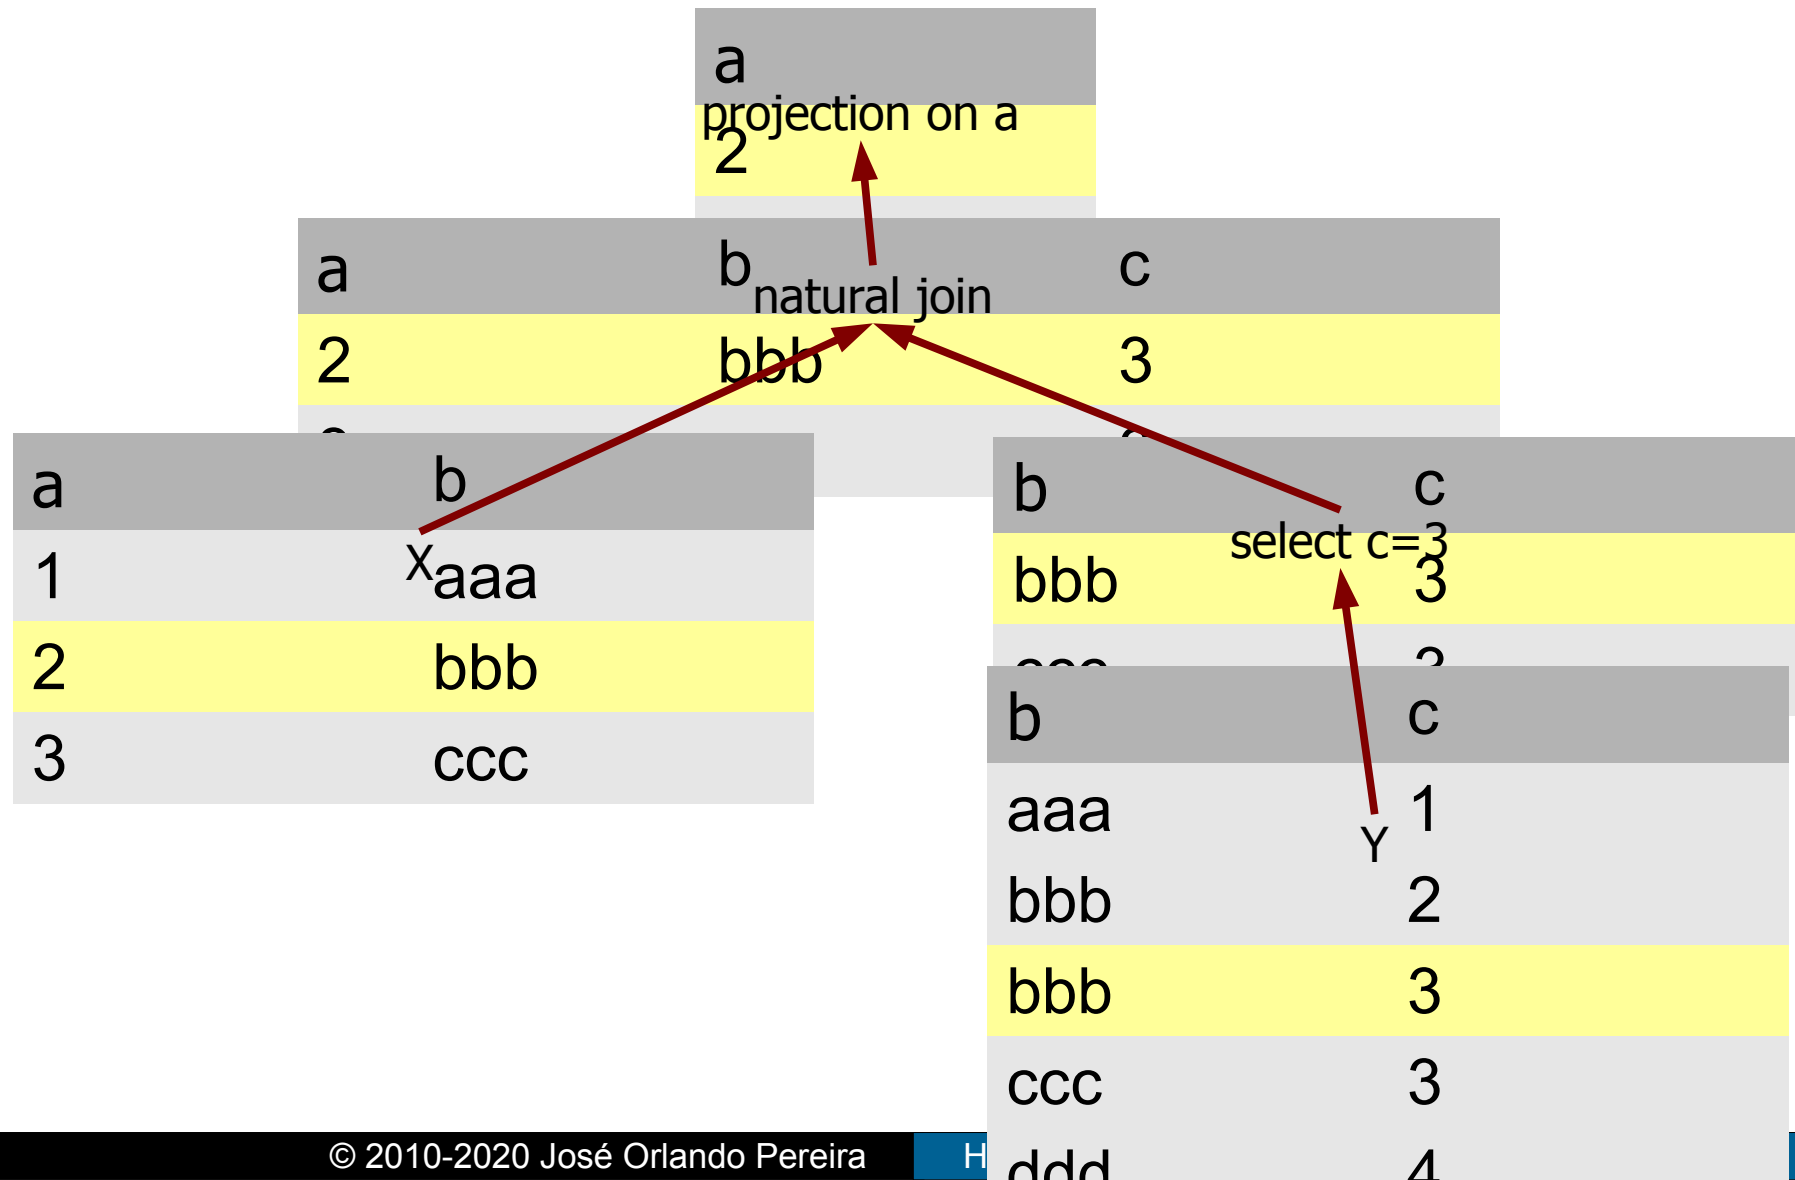


# Execution with iteration






# Execution with iteration



# Execution with iteration



# Materialization vs Iteration

- Iteration avoids caching entire relations
- Materialization avoid reading records more than once
- Can mix both:
  - A materialization operator obtains all records upon first invocation of open
  - Returns records from cached copy on iteration

# Roadmap

- What physical operators exist for each logical operation?
- How are physical operators selected?

# One-pass, record-at-a-time

- Operators:
  - Sequential scan
  - Selection
  - Projection
- Memory requirements:
  - No more than one record required
  - Always possible

# One-pass, full relation, unary

- Duplicate elimination:
  - Cache unique records
  - "select distinct \* from X;"
- Grouping and aggregation:
  - Cache groups
  - "select count(\*) from X group by b;"
- Sorting:
  - Cache all records and sort in memory
  - "select \* from X order by b;"


# One-pass, full relation, binary

- Union, difference, intersection, product, join:
  - Read and cache the smallest relation
  - Organize for fast look-up (e.g. hash)
  - Read and operate on each record from the largest relation

# One-pass, full relation, binary

- Load smaller table into memory and add search structure:

a	b	c
2	bbb	3
3	ccc	3



a	b
1	aaa
2	bbb
3	ccc

	c
bbb	3
ccc	3




# One-pass, full relation, binary

- Test each record from the largest relation:

a	b	c
2	bbb	3
3	ccc	3

a	b
1	aaa
2	bbb
3	ccc



Not found

c
bbb
ccc


# One-pass, full relation, binary

- Test each record from the largest relation:

a	b	c
2	bbb	3
3	ccc	3

a	b
1	aaa
2	bbb
3	ccc

c
bbb
ccc



# Nested-loop join (NLJ)

a	b
1	aaa
2	bbb
3	ccc

b	c
bbb	3
ccc	3
b	c

a	b
1	aaa
2	bbb
3	ccc

b	c
bbb	3
ccc	3
b	c
bbb	3
ccc	3

...

# Nested-loop join (NLJ)

- Memory requirements:
  - One record from each relation
- Operations:
  - If outer loop has  $N$  records
  - Reads inner relation  $N$  times

# Block-based NLJ

- Much smarter: Execute NLJ by blocks
- Memory requirements:
  - One block from each relation
- Operations:
  - If outer loop has  $N$  records /  $B$  blocks ( $B \ll N$ )
  - Reads inner relation  $B$  times ( $B \ll N!$ )

# Large relations and sorting

- Algorithms using sorted data are more efficient (e.g. than nested loops)
- How to sort data that does not fit in memory?

# Merge-sort

- Split data in chunks that fit in memory:

a	b
8	...
4	
7	
2	
3	
5	
6	
1	



a	b
8	...
4	
7	
...	

a	b
3	...
5	
6	
1	

# Merge-sort

- Load and sort each of them:

a	b
2	...
4	
7	
8	

a	b
3	...
5	
6	
1	



# Merge-sort

- Load and sort each of them:

a	b
2	...
4	
7	
8	

a	b
1	...
3	
5	
6	

# Merge-sort

- When iterating, select the next record from the fragment with the next key:

a	b
2	...
4	
7	
8	

a	b
1	...
3	
5	
6	

# Merge-sort

- When iterating, select the next record from the fragment with the next key:

a	b
2	...
4	
7	
8	

a	b
1	...
3	
5	
6	

# Merge-sort

- When iterating, select the next record from the fragment with the next key:

a	b
2	...
4	
7	
8	

a	b
1	...
3	
5	
6	

# Two-pass, full relation, unary

- First pass is sorting
- Duplicate elimination:
  - Cache last record
  - "select distinct \* from X;"
- Grouping and aggregation:
  - Cache last group
  - "select count(\*) from X group by b;"

# Two-pass, full relation, binary

- Union, difference, intersection, product, join:
  - Read record R1 from sorted relation T1
  - Read record R2 from sorted relation T2
  - If  $R1 = R2$ :
    - Use R1 and R2
  - If  $R1 < R2$ :
    - R1 does not exist in T2
    - Skip R1
  - If  $R2 < R1$ 
    - ....

# Conclusion

- There are a number of options for executing each query
- Varying performance:
  - Memory requirements
  - Number of iterations
  - Disk accesses
- What is the best one?
- How can it be discovered?