

Foundations of Distributed Systems

José Orlando Pereira

HASLab / Departamento de Informática
Universidade do Minho




2020/2021

Motivation

- What is truly unique about distributed systems?

Case study: Leader election

- Select a unique leader in a distributed system
- Useful for:
 - Coordination
 - Efficiency
 - ...




Abstraction

- No sockets/clients/servers/
byte buffers/threads/...
- Reasoning in terms of:
 - n processes: sequences of discrete computation events
 - n^2 channels: connect send/receive events in pairs of processes



Case study: Simple algorithm


- Each process trying to be the leader sends its network address to all others
- Each process considers the process with the highest address to be the leader



In practice

- Tight timeouts are dangerous:

- E.g., proportional to mean delay
- Means low coverage




- Large timeouts are not useful:

- E.g., proportional to high percentile
- Taking advantage of time causes a very large performance penalty

In practice

- Solutions that do not use time might have better performance:
 - Run time proportional to mean delay
 - Even if more message exchanges are necessary




Asynchronous system model

- Assume no global time reference
- Assume no bounds on:
 - clock drift
 - processing time
 - message passing time
- Can we still solve the problem?

Example: Approach


- Start with a synchronous reliable fully connected network
- Relax the system model:
 - Unbounded message loss
 - Large/unknown graph diameter
 - Dynamic graph
- Example: Leader election

Example: Leader election



Summary

- A system model is a set of assumptions:
 - “what we believe about the world”
- A system model is an abstraction
- An algorithm solves a problem in a system model:



Summary



- Asynchronous system model abstracts:
 - Heterogeneity
 - Dynamics
 - **Uncertainty**
- Much simpler than handling them explicitly
- Leads to widely applicable solutions

Summary

- What is truly unique about distributed systems?
 - **Uncertainty!**
- Challenge:
 - Perform a computation in spite of uncertainty