


Goals

- No global omniscient observer
- How to observe/reason about the system from the inside?




Example: Distributed deadlock

- Remote invocation
- All processes request and reply to invocations
- A mutex is held while invoking remotely or handling remote invocations
- Distributed deadlock possible when multiple processes invoke each other

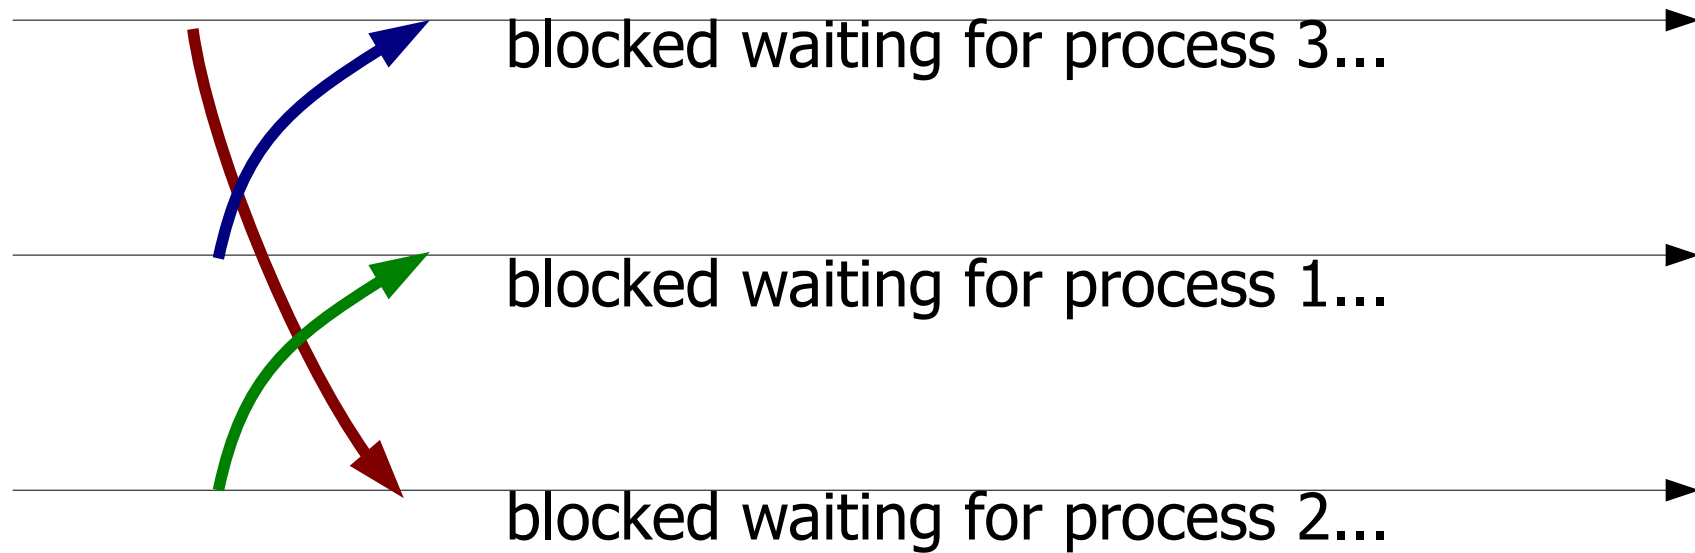
Example: Distributed deadlock

- Deadlock-free run:




Example: Distributed deadlock

- Distributed deadlock:




Example: Distributed deadlock

- Instant observation is impossible:




Example: Distributed deadlock

- Deadlock detection with a “wait for” graph:




Example: Distributed deadlock

- A more complex deadlock-free run:




Example: Distributed deadlock

- A deadlock-free WFG:



Example: Distributed deadlock

- A WFG with a ghost deadlock:




Global Property Evaluation

- Similar problems:
 - Distributed garbage collection
 - Distributed conversation threads
 - ...
- Can it be solved in an asynchronous system?
- Methods that can be used? Relative cost?

Passive monitor process


- Report all events to monitor:



First try: Synchronous system

- Global clock, δ upper bound on message delay
- Tag events with real time
- Consider events only up to $t - \delta$
 - With synchronous rounds, this means using messages from the previous round!

First try: Synchronous system




Clock properties

- What properties of a real-time clock make this approach correct?

Definition: Causality

- Events i and j are causally related ($i \rightarrow j$) iff:
 - i precedes j in some process p
 - for some m , $i = \text{send}(m)$ and $j = \text{receive}(m)$
 - for some k , $i \rightarrow k$ and $k \rightarrow j$ (transitivity)
- Events i and j are concurrent ($i \parallel j$) iff neither $i \rightarrow j$ or $j \rightarrow i$

Causality




Clock properties

- $RC(i)$ the time at which i happened
- If $i \rightarrow j$ then $RC(i) < RC(j)$
- For some event j :
 - When we are sure that there is no unknown i such that $RC(i) < RC(j)$
 - Then there is no i such that $i \rightarrow j$
- Can we build a logical clock with the same property?

Scalar logical clock


- Local events: increment counter
- Send events: increment and then tag with counter
- Receive events: update local counter to maximum and then increment



Second try: Logical clock

- Use scalar logical clock
- Use FIFO channels
- Consider events only up to the minimum of maximum tags

Second try: Logical clock



Scalar clocks

- Synchronous system (RC):
 - Delay δ to consistency
- Asynchronous system (LC):
 - Possible unbounded delay to consistency
 - Blocks if some process stops sending messages