



Vector clock

- Local event at i : increment counter i
- Send event at i : increment counter i and tag with vector
- Receive event at i : update each counter to maximum and increment counter i



Third try: Vector clock



Causal delivery to monitor

- The monitor considers events as follows:
 - With local vector $I[\dots]$
 - For some event $r[\dots]$ from i
 - Ignore it until:
 - $I[i] + 1 = r[i]$
 - For all $j \neq i$: $r[j] \leq I[j]$

Causal delivery in a group

- Broadcast messages in a group
 - Same as “All processes are monitoring send events”
- Increment local counter only on send
 - Not on receive
 - Not on internal events

Summary

- With scalar clocks:
 - Consistent observation by ignoring some messages
 - Blocks if one process stops sending messages
- With vector clocks and causal delivery:
 - Consistent observation whenever a message is delivered
 - Blocking can be avoided by forwarding past messages

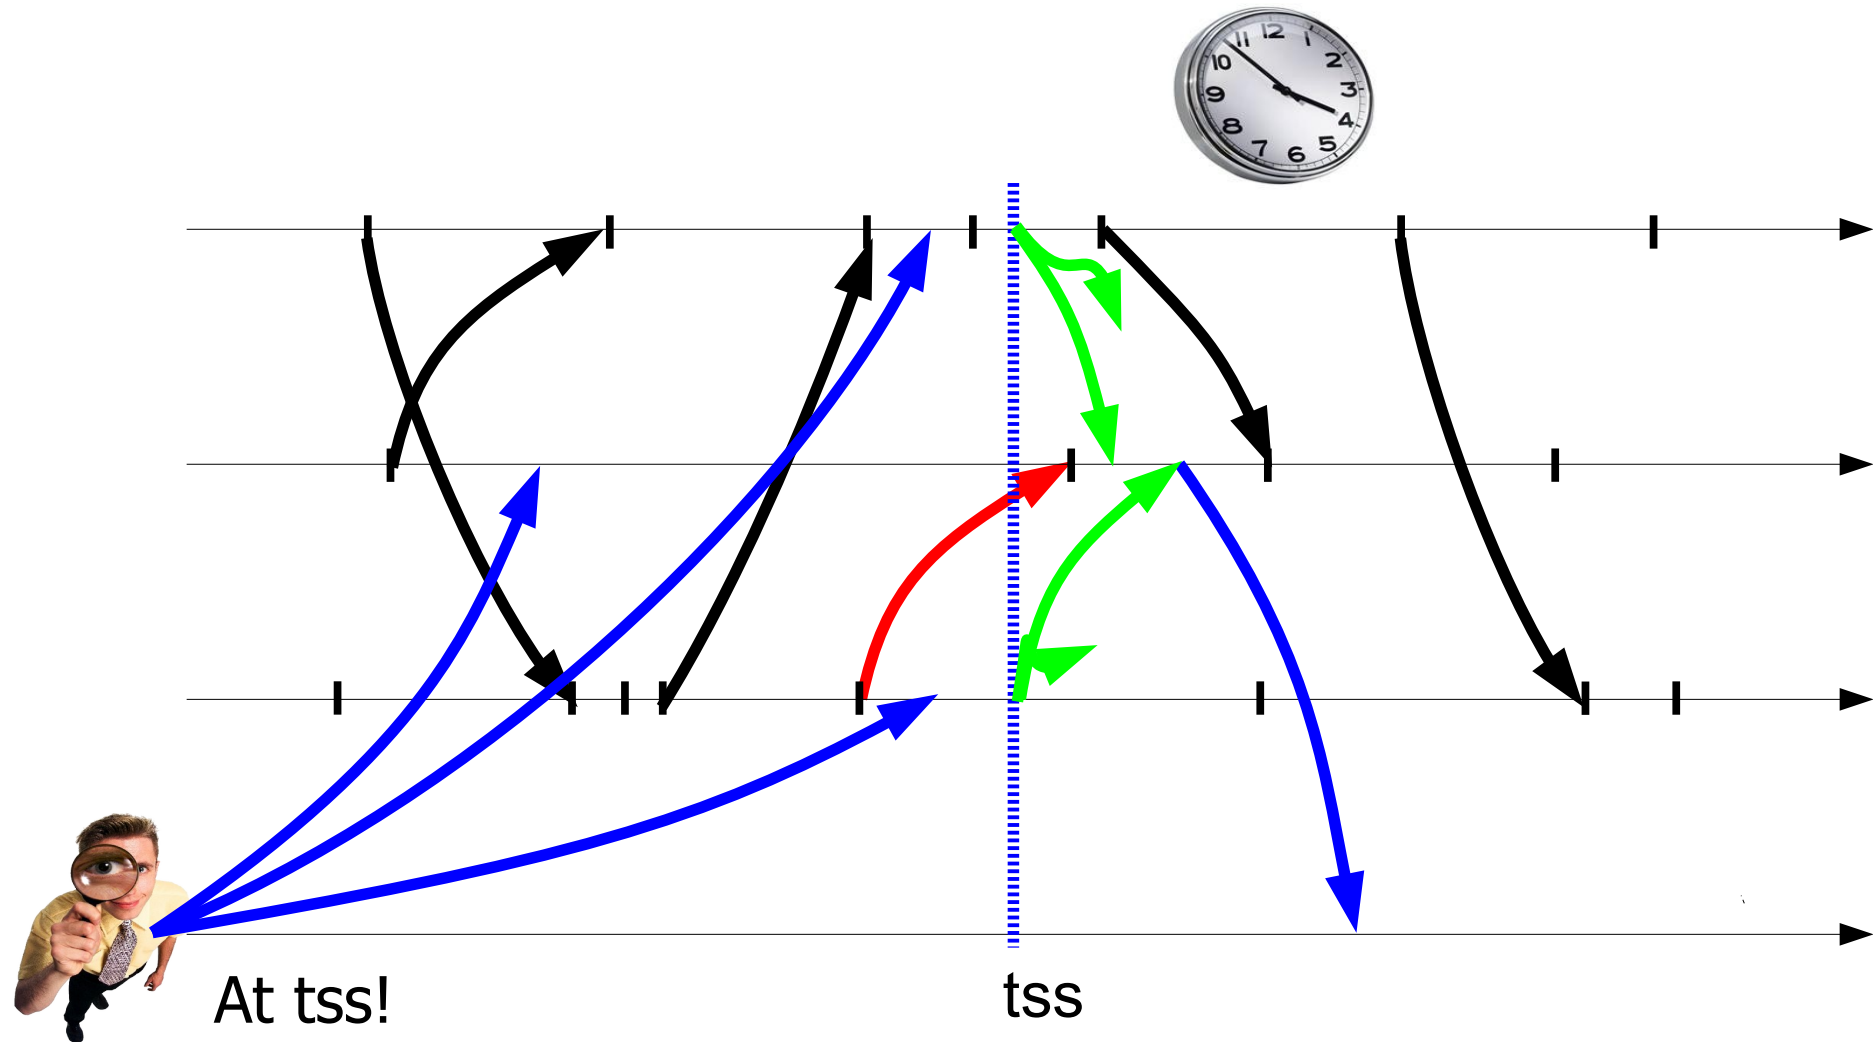
No reporting to monitor process

- Reporting all events to a monitor causes a large overhead
- Can a query be issued at some point in time?


Fourth try: No reporting, synchronous

- Monitor broadcasts tss in the future
- At tss, each process:
 - Records state
 - Sends messages to all others
 - Starts recording messages until receiving a message with $RC > tss$
- After stopping, sends all data to monitor

Fourth try: No reporting, synchronous




Fifth try: No reporting, logical clock



Chandy and Lamport

- Send a "Snapshot" message to some process
- Upon receiving for the first time:
 - Records state
 - Relays "Snapshot" to all others
 - Starts recording on each channel until receiving "Snapshot"
- Send all data to monitor

Chandy and Lamport




Summary

- Can observe the distributed system by:
 - Using scalar logical clocks
 - Using vectorial clocks and causal delivery
 - Using Chandy-Lamport algorithm to collect a discrete snapshot


Cuts and consistency

- A cut is the union of prefixes of process history
- A consistent cut includes all causal predecessors of all events in the cut
- Intuitive methods:
 - If a cut is an instant, there are no messages from the future
 - In the diagram, no arrows enter the cut
 - All events in the frontier are concurrent

Consistent cuts


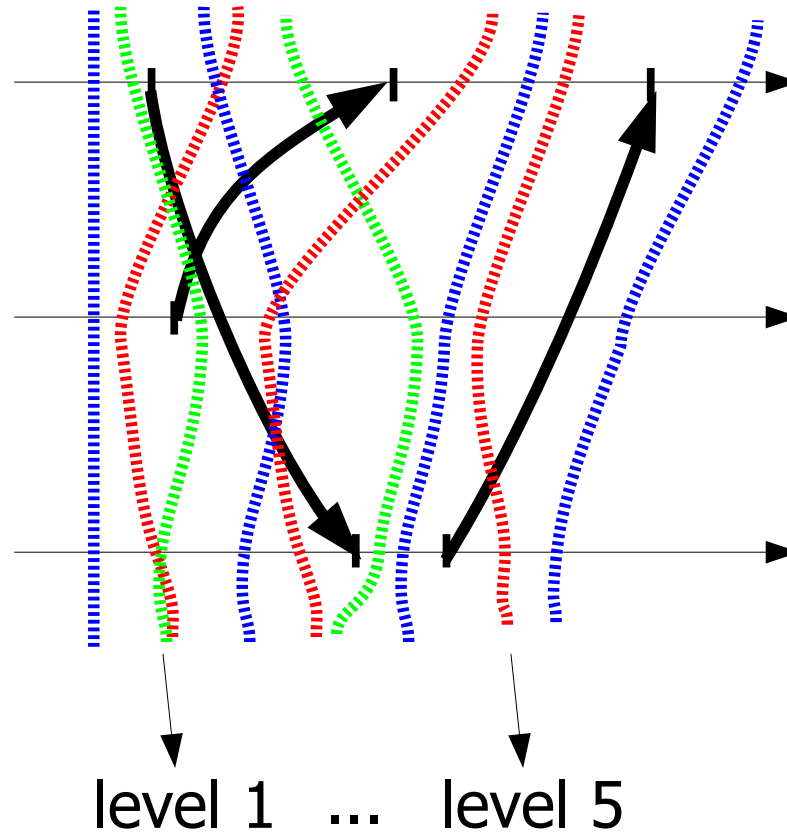


Consistent cuts and global states



- Notation Σ_{625} means state after 6 events in first process, 2 events in second, ...
- This is a *level 13* state (after 6+2+5 events)


State lattice



(*) Not all possible configurations are represented!

Consistent global states



- Includes the true sequence of states in the system
- An observer within the system cannot deny any of the possible paths



Stable predicates

- Once true, always true
- Examples:
 - Deadlock detection
 - Termination
 - Loss of token
 - Garbage collection
- Can be evaluated periodically on snapshots

Stable predicates




Non-stable predicates

- Examples:
 - Total size of queues in the system
 - Number of messages in transit
 - Amount of memory used
- Can be detected by full monitoring of all (relevant) events


Non-stable predicates

- True in a subset of observable states
- Some are possibly true: an observer in the system cannot deny having been true
- The predicate does not hold on some paths



Non-stable predicates

- True in a subset of observable states
- Some are definitely true: an observer in the system is sure of having been true
- The predicate holds on all possible paths



Non-stable predicates

- Start with level $n=1$
- Loop while more states can be found:
 - Generate all level n states (by selecting all messages that can be accepted in state $n-1$)
 - If true in all of these states:
 - return **DEFINITELY TRUE**
 - If true in any of these states:
 - return **POSSIBLY TRUE**
 - Increment n
- return **FALSE**

Summary

- An instant in the asynchronous system model is defined by a consistent cut
 - Safe opportunities to observe the system
- Considering each consistent cut:
 - We cannot deny that the system might have been in that state
 - We can only assert that the system has been in that state if there are no concurrent cuts