


# Distributed Mutual Exclusion

- Implement lock()/unlock() primitives in a distributed system
- Properties:
  - No two processes concurrently in the critical section
  - Some willing process eventually enters the critical section (weak fairness)
  - All willing processes eventually enter the critical section (strong fairness)

# Centralized


coordinator  
manages  
the queue



- At least one round-trip to enter
- Coordinator handles all messages (bottleneck)


# Ring

No explicit queue!



- $N/2$  hops to enter
- Distributed load, but not quiescent

# Physical time




- $\delta$  delay to enter
- Distributed load, but synchronous

# Physical time

- Algorithm:
  - At  $t$ , consider all requests up to  $t - \delta$
  - Order by timestamp, break ties by process id
- $\delta$  delay to enter
- Distributed load, but synchronous

# Logical time



# Logical time

- Algorithm:
  - $ri[j]$  latest timestamp from  $j$  at  $i$
  - Consider requests with  $t \leq \min(ri[j], \text{for all } j)$
  - Order by timestamp, break ties by process id
- 1 hop to enter, if processes send messages frequently
- Distributed load

# Replicated state machine

- Note that all processes keep copies of the queue waiting for the lock
- Can be generalized for any data structure and deterministic computation
  - Replicated state machine